

# Sensing, Predicting, and Planning: Robotic Expeditionary Science in Spatiotemporal Fields

*by*

Victoria Lynn Preston

B.S., Olin College of Engineering (2016)

M.S., Massachusetts Institute of Technology (2019)

Submitted to the Department of Aeronautics and Astronautics  
and the Joint Program in Applied Ocean Science and Engineering  
in partial fulfillment of the requirements for the degree of

## Doctor of Philosophy

*at the*

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

*and*

WOODS HOLE OCEANOGRAPHIC INSTITUTION

February 2023

© 2023 Victoria Lynn Preston. All rights reserved.

The author hereby grants to MIT permission to reproduce and to distribute publicly paper  
and electronic copies of this thesis document in whole or in part in any medium now  
known or hereafter created.

Author .....

Department of Aeronautics and Astronautics, MIT

Applied Ocean Science and Engineering, WHOI

December 15, 2022

Certified by .....

Nicholas Roy

Professor of Aeronautics and Astronautics, MIT

Thesis Supervisor

Certified by .....

Anna Michel

Associate Scientist with Tenure, Applied Ocean Physics and Engineering, WHOI

Thesis Supervisor

Accepted by .....

Jonathan How

Professor of Aeronautics and Astronautics, MIT

Chair, Graduate Committee

Accepted by .....

David Ralston

Associate Scientist with Tenure, Applied Ocean Physics and Engineering, WHOI

Chair, Joint Committee on Applied Ocean Science and Engineering



# Sensing, Predicting, and Planning: Robotic Expeditionary Science in Spatiotemporal Fields

*by*

Victoria Lynn Preston

Submitted to the Department of Aeronautics and Astronautics and the Joint Program in Applied Ocean Science and Engineering on December 15, 2022, in partial fulfillment of the requirements for the degree of Doctor of Philosophy.

## Abstract

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nam quis neque et erat laoreet finibus at ac leo. Curabitur pellentesque, diam quis dignissim finibus, enim dui feugiat leo, nec porttitor sapien mi ac felis. Nam aliquam pretium nibh, quis dapibus dolor gravida sit amet. Cras porttitor dui quis elementum pulvinar. Nulla id pulvinar massa. Nullam ut diam non lorem venenatis faucibus. Vivamus lacus ante, pellentesque vitae nisl sit amet, bibendum facilisis purus.

Thesis Supervisor: Nicholas Roy

Title: Professor of Aeronautics and Astronautics, MIT

Thesis Supervisor: Anna Michel

Title: Associate Scientist with Tenure, Applied Ocean Physics and Engineering, WHOI



# Acknowledgments

Thank you, everyone.



# Contents

<b>1</b>	<b>Introduction</b>	<b>13</b>
<b>2</b>	<b>Problem Setting</b>	<b>15</b>
2.1	Science Background . . . . .	17
2.2	Robotics Background . . . . .	17
<b>3</b>	<b>Sensing: Interpreting Observations of Spatiotemporal Phenomena</b>	<b>19</b>
3.1	Sensing Hydrothermal Fluids from Afar . . . . .	19
3.2	Interpreting Point Observations of Hydrothermal Fluids . . . . .	19
3.3	Leveraging Sensors of Opportunity . . . . .	19
3.4	Summary . . . . .	19
<b>4</b>	<b>Predicting: Representing Belief and Simulating Uncertain Dynamics</b>	<b>21</b>
4.1	General Models: GPs . . . . .	21
4.2	Embedded Physical Models: PHUMES . . . . .	21
4.3	Physically-Informed Learned Models: PIKL . . . . .	21
4.4	Performance in Hydrothermal Domains . . . . .	21
4.5	Summary . . . . .	21
<b>5</b>	<b>Planning: Adaptive Sampling in Spatiotemporal Fields</b>	<b>23</b>
5.1	Underway Autonomy: PLUMES . . . . .	23
5.2	Deployment-by-Deployment Autonomy: PHORTEX . . . . .	23
5.3	Field Results in GB . . . . .	23

5.4	Summary . . . . .	23
<b>6</b>	<b>Discussion</b>	<b>25</b>
6.1	Generalizing to Other Domains . . . . .	25
6.2	Open Challenges . . . . .	25
6.3	Embedding Autonomy Engineers into Science Teams . . . . .	25
6.4	Robotics in Expeditionary Science . . . . .	25
<b>7</b>	<b>Looking Forward</b>	<b>27</b>
<b>A</b>	<b>Tables</b>	<b>29</b>



# List of Figures



# List of Tables

A.1	Armadillos . . . . .	29
-----	----------------------	----



# Chapter 1

## Introduction

The environmental sciences are a multidisciplinary endeavor to understand the Earth, its ecosystems, and its processes, for which *in situ* observational studies or *expeditions* serve as the foundation on which scientific discovery and model development is predicated. Robots are uniquely well-positioned to advance long-term monitoring of and exploration in meso-scale planetary environments through autonomous expeditions. By virtue of their form, robots can be used in extreme places (e.g., deep sea), dangerous scenarios (e.g., edge of calving ice sheets), or long-term missions (e.g., Mars exploration). Increasingly, robotic platforms are being developed for scientific expeditions, but their autonomous capabilities are typically limited to predetermined hand-designed trajectories (e.g., uniform coverage lawnmowers [Camilli et al., 2010]). This significantly restricts the class of phenomena that can be effectively studied by expeditionary scientific robots. For instance, spatiotemporal distributions—deep sea hydrothermal plumes, algal blooms, weather cells—can be severely under-sampled or missed using these strategies [Flaspohler et al., 2019]. Given the ubiquity of these spatiotemporal phenomena and the cost of scientific field operations, it is critical to improve the efficacy of robots as autonomous scientific tools.

This thesis represents...



# Chapter 2

## Problem Setting

To collect useful samples of a spatiotemporal field using a robotic platform is to pose a *sequential decision-making* problem. In this setting, we assume that the measurements that can be collected are partial observations of the unknown spatiotemporal environment, and the actions the robot can take in sequence are operationally constrained. We can formally state this problem as a partially observable Markov decision process (POMDP). Let  $\Pi(\cdot)$  denote the space of probability distributions over the argument. A finite horizon POMDP can be represented as tuple:  $(\mathcal{S}, \mathcal{A}, T, R, \mathcal{Z}, O, b_0, H, \gamma)$ , where  $\mathcal{S}$  are the states,  $\mathcal{A}$  are the actions, and  $\mathcal{Z}$  are the observations. At planning iteration  $t$ , the agent selects an action  $a \in \mathcal{A}$  and the transition function  $T : \mathcal{S} \times \mathcal{A} \rightarrow \Pi(\mathcal{S})$  defines the probability of transitioning between states in the world, given the current state  $s$  and control action  $a$ . The transition function governs both how the state of the robot will evolve, given a chosen action, and the potentially stochastic evolution of the underlying spatiotemporal environment. After the state transition, the agent receives an observation according to the observation function  $O : \mathcal{S} \times \mathcal{A} \rightarrow \Pi(\mathcal{Z})$ , which defines the probability of receiving an observation, given the current state  $s$  and previous control action  $a$ . The reward function  $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  serves as a specification of the task, assigning the states of the world that are useful for a given scientific objective high reward and others low reward. A POMDP is initialized with belief  $b_0 \in \Pi(\mathcal{S})$  — an initial probability distribution over state — and plans over horizon  $H \in \mathbb{Z}^+$  with discount factor

$\gamma \in [0, 1]$ .

As the robot moves through the world, it selects actions and receives observations. Since the state of the world is not directly observable in a POMDP, the robot maintains a probability distribution over possible states (i.e., belief) and must update this distribution each time it takes an action and receives an observation. Given the transition and observation models, the belief can be updated directly using Bayes rule using a Bayes filter [Särkkä, 2013]:

$$\tau(b_{t-1}, a_{t-1}, z_t) = b_t \triangleq \Pi(S_t \mid a_0, z_0, \dots, a_{t-1}, z_{t-1}, z_t) \quad (2.1)$$

$$= \Pi(S_t \mid b_{t-1}, a_{t-1}, z_t) \quad (2.2)$$

$$= \frac{\int_{s \in \mathcal{S}} O(s, a_{t-1}, z_t) T(s, a_{t-1}, s') b_{t-1}(s')}{\Pi(z_t \mid b_{t-1}, a_{t-1})} \quad (2.3)$$

where  $\tau(b, a, z)$  is the updated belief after taking control action  $a$  and receiving observation  $z$  (Eq. (2.3)). Unfortunately, Eq. (2.3) is intractable to compute directly and an approximate Bayesian inference procedure is required to represent the belief (e.g., Kalman filter [Welch et al., 1995], particle filter [Silver and Veness, 2010], or variational methods).

Due to the stochastic, partially observable nature of current and future states, the realized reward in a POMDP is a random variable. Optimal planning is defined as finding a horizon-dependent policy  $\{\pi_t^* : \Pi(\mathcal{S}) \rightarrow \mathcal{A}\}_{t=0}^{H-1}$  that maximizes expected reward:  $\mathbb{E}\left[\sum_{t=0}^{H-1} \gamma^t R(S_t, \pi_t(b_t)) \mid b_0\right]$ , where  $b_t$  is the updated belief at time  $t$ , conditioned on the history of actions and observations. The recursively defined horizon- $h$  optimal value function  $V_h^*$  quantifies, for any belief  $b$ , the expected cumulative reward of following an optimal policy over the remaining planning iterations:  $V_0^*(b) = \max_{a \in \mathcal{A}} \mathbb{E}_{s \sim b}[R(s, a)]$  and

$$V_h^*(b) = \max_{a \in \mathcal{A}} \mathbb{E}_{s \sim b}[R(s, a)] + \gamma \int_{z \in \mathcal{Z}} \Pi(z \mid b, a) V_{h-1}^*(\tau(b, a, z)) dz \quad h \in [1, H-1], \quad (2.4)$$

The optimal policy at horizon  $h$  is to act greedily according to a one-step look ahead of



the horizon- $h$  value function. However, Eq. (2.4) is intractable for large or continuous state, action, or observation spaces and thus the optimal policy must be approximated. Much of the art of practical decision-making under uncertainty is making well-designed algorithmic and heuristic choices that enable efficient and robust planning algorithms.

## **2.1 Science Background**

## **2.2 Robotics Background**



## Chapter 3

# Sensing: Interpreting Observations of Spatiotemporal Phenomena

### 3.1 Sensing Hydrothermal Fluids from Afar

TRANSECT PAPER!

### 3.2 Interpreting Point Observations of Hydrothermal Fluids

PSEUDO-SENSOR

### 3.3 Leveraging Sensors of Opportunity

PIV, Tiltmeters, Observatories, Now-casts, Forecasts

### 3.4 Summary



## Chapter 4

# Predicting: Representing Belief and Simulating Uncertain Dynamics

4.1 General Models: GPs

4.2 Embedded Physical Models: PHUMES

4.3 Physically-Informed Learned Models: PIKL

4.4 Performance in Hydrothermal Domains

4.5 Summary



## Chapter 5

# Planning: Adaptive Sampling in Spatiotemporal Fields

### 5.1 Underway Autonomy: PLUMES

### 5.2 Deployment-by-Deployment Autonomy: PHOR- TEX

### 5.3 Field Results in GB

### 5.4 Summary





# Chapter 6

## Discussion

6.1 Generalizing to Other Domains

6.2 Open Challenges

6.3 Embedding Autonomy Engineers into Science Teams

6.4 Robotics in Expeditionary Science



# Chapter 7

## Looking Forward

This is the conclusion section.



# Appendix A

## Tables

Table A.1: Armadillos

Armadillos	are
our	friends



# Bibliography

- Camilli, R., Reddy, C. M., Yoerger, D. R., Van Mooy, B. A., Jakuba, M. V., Kinsey, J. C., McIntyre, C. P., Sylva, S. P., and Maloney, J. V. (2010). Tracking hydrocarbon plume transport and biodegradation at deepwater horizon. *Science*, 330(6001):201–204.
- Flaspohler, G., Preston, V., Michel, A. P., Girdhar, Y., and Roy, N. (2019). Information-guided robotic maximum seek-and-sample in partially observable continuous environments. *IEEE Robotics and Automation Letters*, pages 3782–3789.
- Särkkä, S. (2013). *Bayesian filtering and smoothing*. Number 3. Cambridge university press.
- Silver, D. and Veness, J. (2010). Monte-Carlo planning in large POMDPs. In *Proc. Adv. Neural Inf, Process, Syst.*, pages 2164–2172.
- Welch, G., Bishop, G., et al. (1995). An introduction to the kalman filter.